



Predicting Loan Defaults

Team Sparkling

Mengxin Qian, Will Young, April Wang, Jinxin Ma

Outline

- Data & Introduction
- Feature Engineering
- Modeling
- Results and Evaluation
- Lessons Learned




Data & Introduction

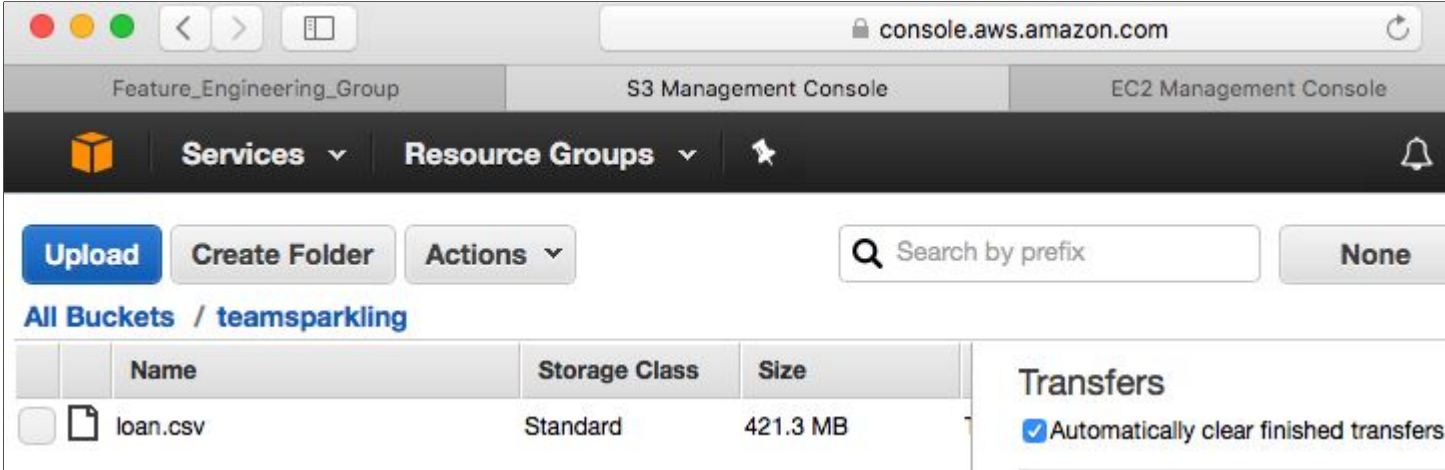


- Analytics goal
 - **Predict whether a customer will pay off the loan**
- Why we chose this data:
 - Interesting and challenging: 73 total features, 880,000 observations
 - Successful predictions increase profit for investors

Data processing



Data in S3 Bucket




The screenshot displays the AWS Management Console interface for an S3 bucket named 'teamsparkling'. The browser address bar shows 'console.aws.amazon.com'. The console has tabs for 'Feature_Engineering_Group', 'S3 Management Console', and 'EC2 Management Console'. The 'S3 Management Console' tab is active. The top navigation bar includes 'Services', 'Resource Groups', and a notification bell. Below the navigation bar, there are buttons for 'Upload', 'Create Folder', and 'Actions'. A search bar labeled 'Search by prefix' and a 'None' button are also present. The main content area shows the bucket path 'All Buckets / teamsparkling'. A table lists the contents of the bucket:

	Name	Storage Class	Size
<input type="checkbox"/>	loan.csv	Standard	421.3 MB

To the right of the table, there is a section titled 'Transfers' with a checkbox labeled 'Automatically clear finished transfers' which is checked.

Create Data Frame & Spark SQL



1. Import Libraries

```
In [1]: from pyspark.sql.types import *
        from pyspark.ml.feature import StringIndexer, OneHotEncoder
        from pyspark.sql.functions import *
```

2. Load Data and Convert to Spark Data Frame

```
In [2]: # load data as dataframe
        loan_df = spark.read.csv("s3n://teamsparkling/loan.csv", header=True)
```

```
In [3]: loan_df.rdd.getNumPartitions()
```

Out[3]: 64

```
loan_df_final.select('paid_flag',
                      'loan_amnt',
                      'funded_amnt',
                      'int_rate',
                      'installment',
                      'annual_inc').show(10)
```

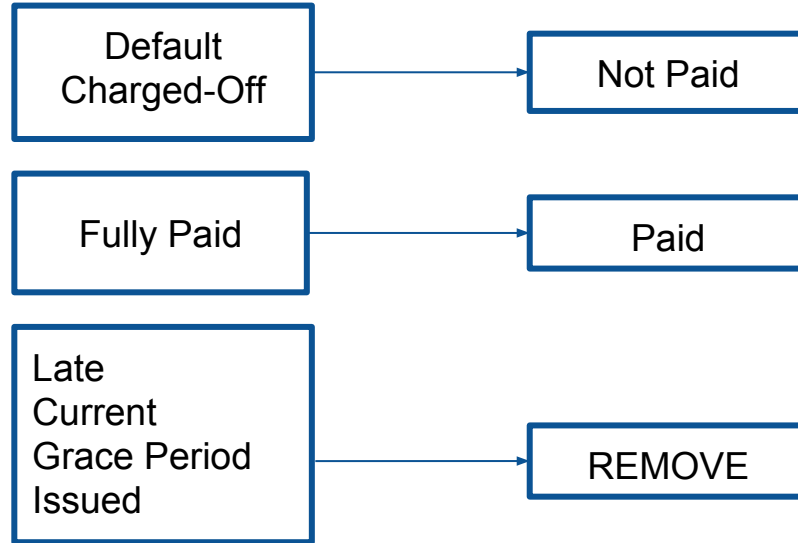
paid_flag	loan_amnt	funded_amnt	int_rate	installment	annual_inc
0	5000.0	5000.0	10.65	162.87	24000.0
1	2500.0	2500.0	15.27	59.83	30000.0
0	2400.0	2400.0	15.96	84.33	12252.0
0	10000.0	10000.0	13.49	339.31	49200.0
0	5000.0	5000.0	7.9	156.46	36000.0
0	3000.0	3000.0	18.64	109.43	48000.0
1	5600.0	5600.0	21.28	152.39	40000.0
1	5375.0	5375.0	12.69	121.45	15000.0
0	6500.0	6500.0	14.65	153.45	72000.0
0	12000.0	12000.0	12.69	402.54	75000.0

only showing top 10 rows

Feature Engineering

- Create response variable
 - Convert data type
 - Bucket feature categories
 - Encode categorical features
 - Handle missing values
-

Feature Engineering - Response Variable



- *Current* does not imply the loan will be paid
 - *Late* does not imply the loan will not be paid
-


Feature Engineering - Data Cleaning

emp_length	emp_len
9 years	null
5 years	1
1 year	6
n/a	3
2 years	5
7 years	9
8 years	4
4 years	8
6 years	7
3 years	10
10+ years	2
< 1 year	

```
import re
def convert_to_int(s):
    s = re.sub('\D', '', s) #remove any non-digital character
    #\d matches any digital, #\D matches any non-digital
    try:
        return s
    except ValueError:
        return None

emp_to_num = udf(convert_to_int)
loan_df4 = loan_df3.withColumn('emp_len', emp_to_num('emp_length').cast('integer'))
```


Feature Engineering - Categorical Features



Feature Engineering - Missing Values

3.3.7 deal with missing values

```
# Filling NA values with 0  
loan_df8 = loan_df8.fillna(0.0, ['tot_coll_amt', 'tot_cur_bal', 'total_rev_hi_lim'])  
# Drop rows with NA  
loan_df9 = loan_df8.dropna()
```

- Filling missing value with 0
 - Drop rows with missing value
-

Logistic Regression & CV parameter tuning

Secure | https://ec2-35-167-233-46.us-west-2.compute.amazonaws.com:8888/notebooks/Feature_Engineering_Group.ipynb

Feature_Engineering_Group (autosaved)

Edit View Insert Cell Kernel Widgets Help

⌘ ⌘ ⌘ ⬆ ⬇ ⏪ ⏴ ⏵ ⏩ Code CellToolbar 📁 📄 📺

4.3 Logistic Regression

```
2]: lr = LogisticRegression()
    evaluator = BinaryClassificationEvaluator()

    cv = CrossValidator().setEstimator(lr).setEvaluator(evaluator).setNumFolds(5)
    paramGrid = ParamGridBuilder().addGrid(lr.maxIter,[100]).addGrid(lr.regParam,[0.1,0.01,0.001,1]).build()

    cv.setEstimatorParamMaps(paramGrid)
    cvmodel = cv.fit(data_train)

4]: predictions = cvmodel.bestModel.transform(data_test)
    score = evaluator.evaluate(cvmodel.bestModel.transform(data_test))
    print evaluator.getMetricName()+" "+str(score)
```

areaUnderROC:0.703155187142

Random Forest & CV parameter tuning

Not Secure https://ec2-35-167-233-46.us-west-2.compute.amazonaws.com:8888/notebooks/Feature_Engineering_Group.ipynb

Jupyter Feature_Engineering_Group (autosaved)

Edit View Insert Cell Kernel Widgets Help


4.2 Random Forest

```
33]: # Fit random forest
      rf = RandomForestClassifier()
      evaluator_rf = BinaryClassificationEvaluator()
      cv = CrossValidator().setEstimator(rf).setEvaluator(evaluator_rf).setNumFolds(5)
      paramGrid = ParamGridBuilder().addGrid(rf.numTrees,[31,51]).addGrid(rf.maxDepth,[5,15]).build()

      cv.setEstimatorParamMaps(paramGrid)
      cvmodel_rf = cv.fit(data_train)
```

Evaluation Metrics


- areaUnderROC : 0.71
- Confusion matrix
 - Precision
 - Recall
 - Accuracy



Initial Test Results

	Random Forest
Precision	0.46
Recall	0.05
Accuracy	0.82

Balancing TRAIN Data



Test Results

	Random Forest (80/20 data)	Logistic Regression Model (50/50 data)	Random Forest Model (50/50 data)	Improvement
Precision	0.46	0.29	0.28	-37.23%
Recall	0.05	0.66	0.65	1215.09%
Accuracy	0.82	0.64	0.64	-21.48%

Test Results

```
metrics = MulticlassMetrics(sc.parallelize(predictionAndLabels))
print metrics.confusionMatrix().toArray()
print metrics.precision(1)
print metrics.recall(1)
print metrics.accuracy
```

```
[[ 25876.  14500.]
 [  3066.   5887.]]
0.288762446657
0.657544956998
0.64390115348
```

	Predicted Paid	Predicted Not Paid	Metrics
Actual Paid	26,000 (53%)	14,500 (30%)	Precision = 0.28
Actual Not Paid	3,000 (5%)	6,000 (12%)	Recall = 0.66

Lessons Learned

- Unbalanced training set can be problematic
 - Recall is crucial in default detection
 - Feature importance is hard to handle in ML
 - Overhead issue with too many partitions
-

Thank you!

